

Name: Gandhi Meet Vipul  
PRN No: 2020BTECS00112

## High Performance Computing Lab

### Practical No. 8

**Title of practical:** Implementation of Vector-Vector addition & N-Body Simulator using CUDA C

#### Problem Statement 1:

Implement Vector-Vector addition using CUDA C. State and justify the speedup using different size of threads and blocks.

#### Screenshots:

a. <<<1,1>>>

##### CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
89.1	2070917912	1	2070917912.0	2070917912	2070917912	cudaDeviceSynchronize
10.4	242855025	3	80951675.0	30158	242733210	cudaMallocManaged
0.5	11048723	3	3682907.7	3633039	3774470	cudaFree
0.0	107264	1	107264.0	107264	107264	cudaLaunchKernel

##### CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	2070966675	1	2070966675.0	2070966675	2070966675	addVectorsInto(float*, float*, float*, int)

b. <<<1,64>>>

##### CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
49.8	225530532	3	75176844.0	17106	225468487	cudaMallocManaged
47.7	215892547	1	215892547.0	215892547	215892547	cudaDeviceSynchronize
2.5	11114074	3	3704691.3	3642477	3796793	cudaFree
0.0	30477	1	30477.0	30477	30477	cudaLaunchKernel

##### CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	215879998	1	215879998.0	215879998	215879998	addVectorsInto(float*, float*, float*, int)

c. <<<1,128>>>

##### CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
59.6	225311780	3	75103926.7	24017	225205463	cudaMallocManaged
37.1	140300558	1	140300558.0	140300558	140300558	cudaDeviceSynchronize
3.2	12118062	3	4039354.0	3983222	4148977	cudaFree
0.0	49955	1	49955.0	49955	49955	cudaLaunchKernel

##### CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	140285660	1	140285660.0	140285660	140285660	addVectorsInto(float*, float*, float*, int)

d. <<<2,64>>>

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
54.7	223015874	3	74338624.7	37099	222849819	cudaMallocManaged
40.5	165007654	1	165007654.0	165007654	165007654	cudaDeviceSynchronize
4.7	19290938	3	6430312.7	6342114	6601342	cudaFree
0.0	163465	1	163465.0	163465	163465	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	165077880	1	165077880.0	165077880	165077880	addVectorsInto(float*, float*, float*, int)

e. <<<2,128>>>

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
64.6	228493155	3	76164385.0	36231	228357618	cudaMallocManaged
30.2	106741694	1	106741694.0	106741694	106741694	cudaDeviceSynchronize
5.2	18491462	3	6163820.7	6061214	6253149	cudaFree
0.0	41798	1	41798.0	41798	41798	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	106726488	1	106726488.0	106726488	106726488	addVectorsInto(float*, float*, float*, int)

f. <<<4,64>>>


CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
70.0	297804934	3	99268311.3	18898	297725683	cudaMallocManaged
25.8	109763552	1	109763552.0	109763552	109763552	cudaDeviceSynchronize
4.2	17874379	3	5958126.3	5917572	6020290	cudaFree
0.0	32427	1	32427.0	32427	32427	cudaLaunchKernel

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	109751222	1	109751222.0	109751222	109751222	addVectorsInto(float*, float*, float*, int)

Output:



Threads \ Blocks	Threads		
	64	128	256
1	0.215	0.14	0.119
2	0.165	0.106	0.069
4	0.109	0.09	0.066

## Problem Statement 2:

Implement N-Body Simulator using CUDA C. State and justify the speedup using different size of threads and blocks.

### Screenshots:

```
In [17]: !nvcc -std=c++11 -o nbody 09-nbody/01-nbody.cu
```

It is highly recommended you use the profiler to assist your work. Execute the following cell to generate a report file:

```
In [18]: !nsys profile --stats=true --force-overwrite=true -o nbody-report ./nbody
```

```
Warning: LBR backtrace method is not supported on this platform. DWARF backtrace method will be used.
WARNING: The command line includes a target application therefore the CPU context-switch scope has been set to process-tree.
Collecting data...
21.620 Billion Interactions / second
Processing events...
Saving temporary "/tmp/nsys-report-653b-3688-e94d-efea.qdstrm" file to disk...
```

```
Creating final output files...
Processing [=====100%]
Saved report file to "/tmp/nsys-report-653b-3688-e94d-efea.qdrep"
Exporting 1109 events: [=====100%]

Exported successfully to
/tmp/nsys-report-653b-3688-e94d-efea.sqlite
```

CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
97.3	294102449	1	294102449.0	294102449	294102449	cudaMalloc
2.5	7589096	20	379454.8	3461	753403	cudaDeviceSynchronize
0.1	183892	1	183892.0	183892	183892	cudaFree
0.1	155654	20	7782.7	3744	28941	cudaLaunchKernel
0.0	105842	2	52921.0	44477	61365	cudaMemcpy

CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
99.4	7477098	10	747709.8	746689	748449	bodyForce(Body*, float, int)
0.6	48320	10	4832.0	4704	5376	integratePosition(Body*, float, int)

```
In [19]: from assessment import run_assessment
```

Execute the following cell to run and assess nbody :

```
In [20]: run_assessment()
```

```
Running nbody simulator with 4096 bodies
-----

Application should run faster than 0.9s
Your application ran in: 0.2122s
Your application reports 21.942 Billion Interactions / second

Your results are correct

Running nbody simulator with 65536 bodies
-----

Application should run faster than 1.3s
Your application ran in: 0.5410s
Your application reports 114.210 Billion Interactions / second

Your results are correct

Congratulations! You passed the assessment!
See instructions below to generate a certificate, and see if you can accelerate the simulator even more!
```

### Output:

For 4096 bodies: 0.2122 secs

For 65536 bodies: 0.5410 secs